

# Generative Adversarial Nets

Short presentation of the paper  
by Goodfellow et al.

Makito Fredskild Katsume

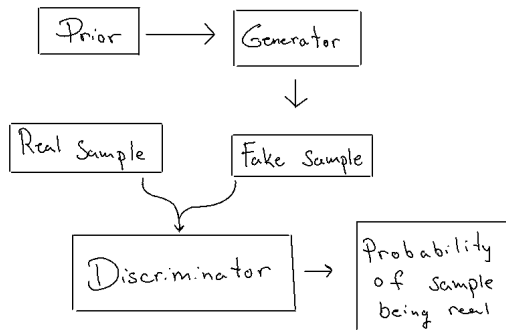
UNIVERSITY OF COPENHAGEN



## What is a generative adversarial network?

- Two models trained simultaneously, the generator and the discriminator.
- A minmax game played between the generator and the discriminator.

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$



## A graphical view

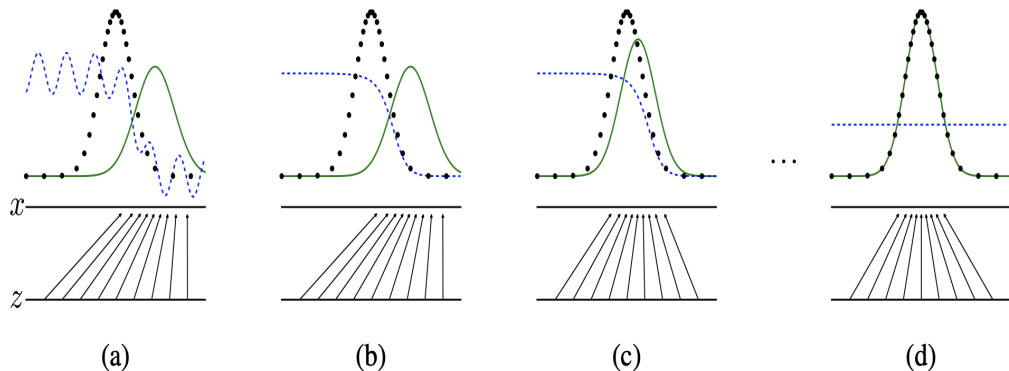


Figure: From "Generative adversarial nets" by Goodfellow et al.[1]

# Algorithm

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

Figure: From "Generative adversarial nets" by Goodfellow et al.[1]

## Theoretical result, the short version

Given:

- Models have infinite capacity and training time
- During each step in the training the discriminator finds the optimum for a given generator

Then:

The distribution of the generator converges to the distribution of the training data[1].

## My application

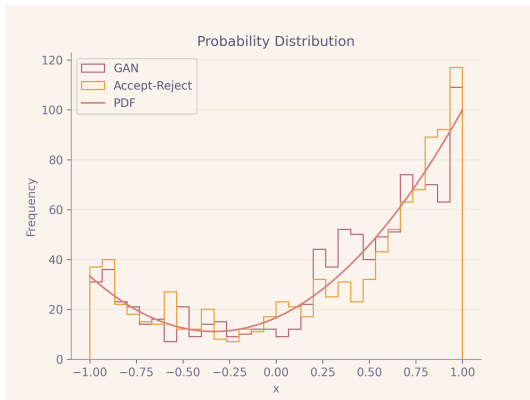
The distribution to be approximated is

$$pdf(x) = C(1 + \alpha x + \beta x^2) \text{ for } x \in [-1, 1]$$


where  $C$  is the normalisation constant,  $\alpha = 2$ , and  $\beta = 3$ .

- Prior space uniformly distributed on  $\mathbf{z} \in [0, 1]$
- Real data is sampled by Accept-Reject method

# Result

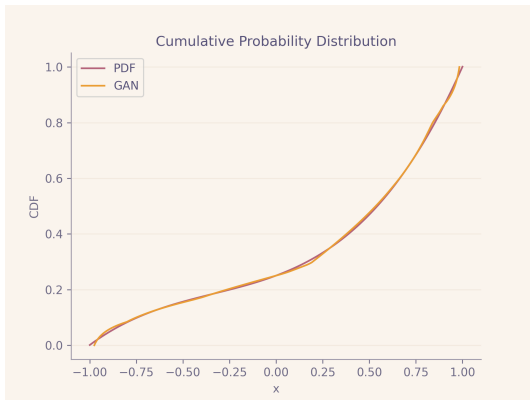


(a) A comparison of samples generated with Accept-Reject method and the GAN model.

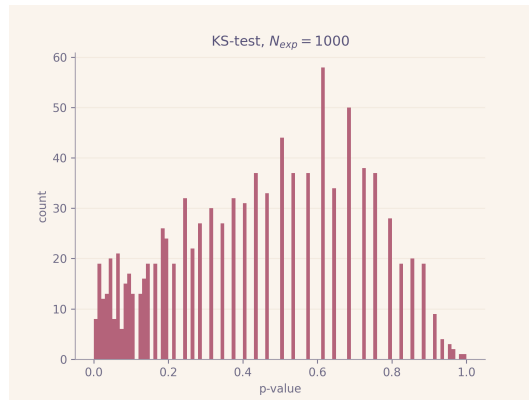


(b) Cumulative distribution of the sample from the left plot.

# Result



(a) Cumulative distribution of the resulting GAN and the training PDF.



(b) KS-Test on 1000 iteration.



## Conclusion

- It kind of works.
- A lot of meta parameters to choose and adjust.
- Needs a lot of labelled data.
- Could be viable in the right use case.

### ***Reference :***

- [1] Ian J. Goodfellow et al. "Generative Adversarial Nets". In: (2014). DOI: <https://arxiv.org/abs/1406.2661>.